UNIVERSITY OF SCIENCE
AND TECHNOLOGY OF HANOI

BACHELOR THESIS

Development of Data Lake Core

Author

Nguyễn Gia Phong

Supervisor

Trần Giang Sơn, PhD

July 12, 2021

Contents

Declaration

I declare that I have composed this thesis in its entirety, as the result of my own work, unless explicitly indicated otherwise via referencing. The presented work has not been submitted for any previous application for a degree or professional qualification.

1

Introduction

1.1 Motivation

Many researchers at University of Science and Technology of Hanoi (USTH) operate with data on a regular basis and often a dataset is studied by multiple researchers from different departments and points in time. Currently the data are organized manually, even on the laboratories' storages, which is prone to duplication and makes data discovery difficult.

A data lake shared among the university's researchers, professors and students will not only save resources but also improve productivity and promote interdisciplinary collaborations. With USTH's goal of growing to be an excellent research university in Việt Nam and in the region [1], building such data lake can be an essential task.

1.2 Background

A *data lake* is a massive repository of multiple types of data in their raw format at scale for a low cost [2]. The data's schema (structure) is defined on read to minimize data modeling and integration costs [2].

For the ease and efficiency of scaling, a microservice architecture could be a good choice. By arranging the data lake as a collection of loosely-coupled services, it becomes possible to scale individual services individually [3]. In this architecture, the *core* microservice is defined as the innermost component, which communicates directly with the storages. The core shall provide an application programming interface (API) for other components to upload, query and extract data.

Append-only storages only allow new data to be appended, whilst ensure the immutability of existing data. As immutable data are thread-safe, they

reduce the complexity of the concurrency model, making it easier to comprehend and reason about [4]. This is particularly useful in large distributed systems with multiple moving parts.

Since the data are immutable, each *content* can be given an identifier (ID), i.e. a content identifier (CID). For end-users, we also introduce a higher level concept: *dataset*, composing of not only the content but also relevant metadata for indexing. Like contents, datasets can also immutable, with changes written as a new revision linking to the previous one.

Append-only storages' operations boil down to two kinds: appending and reading. For the latter, sometimes the data are not wanted in their entirety, but filtered and accumulated. While data of different types usually requires different tools and libraries to query upon, the core API should be providing one single query language for all data types, plus their metadata. In this thesis, such usage is referred to as *query polymorphism.*

1.3   Objectives

The work presented here was done as part of a three-month internship in collaboration with several other students at USTH ICTLab* to build a data lake for a better management of the university's data. The internship focused on the lake's core microservice, which abstracts underlying persistent layers and perform relevant metadata transformation and discovery. It should provide an internal interface to other components for data ingestion, (primitive) query and extraction, as well as carrying out tasks for enhancing the discoverability and usability of the aforementioned datasets.

After the internship period, the resulting codebase shall be maintained by ICTLab and future students, so the work must be designed, implemented and documented in a way that ensures such possibility.

1.4   Expected Outcomes

The intended deliverables of the three-month internship are listed as follows:

- Requirement analysis of the data lake core

- Data lake core's architecture and design

- Core API design and specification

- Implementation and integration with other components

---

*`https://ictlab.usth.edu.vn`

2

Methodology

Development followed the evolutionary prototyping model: a robust prototype is built by improving and adding newly understood features [5]. This often took 7–10 days, with requirements, design, implementation and test suite all being refined.

2.1 Requirement Analysis

In this section, from given context and objectives, we analyzed the expected system for a set of features and derived a list of use cases. Supplementary specifications were also added to elaborate on the nonfuntional requirements.

2.1.1 Use-Case Model

As previously introduced, the most basic functions of the data lake core are content uploading and downloading, along with datasets addition and querying. A more advanced (and rather powerful) use case is content extraction, which allows one to fetch only the interested part of the content, e.g. extracting rows matching a certain predicate from (semi-)structured data. Together with logging, the core's use cases are summarized in figure 2.1.

Upload content   This use case allows other microservices of the data lake to upload a content. Its flow of events is depicted as follows, where error handling is omitted for brevity, since all errors, if occur, replace the normal response.

1. A content is sent to the core microservice.

Figure 2.1: Use-case diagram of the core microservice

2. Core adds the content to the underlying storage and register it to the database (DB).

3. Core responds with the CID of the added content.

Add dataset   This use case let other services add a dataset:

1. A dataset is sent to the core microservice.

2. Core adds the dataset to the underlying DB.

3. Core responds with the ID of the added dataset.

Find datasets   This use case allows other services to find the datasets whose metadata satisfy a given predicate:

1. A predicate is sent to the core microservice.

2. Core runs a query in the underlying DB to find matching datasets.

3. Core responds with a linear collection of metadata, each of which satisfying the given predicate.

Download content   In use case, other services fetch a content from the data lake core.

1. A CID is sent to the core microservice.

2. Core passes the CID to the underlying storage.

3. Core responds with the respective content.

Extract content   This use case allows other services to extract a content's parts satisfying a given predicate:

1. A CID and a predicate is sent to the core microservice.

2. Core iterates the content for matching elements.

3. Core responds with the extracted elements.

Gather logs   This use case let system admins study events occurring in the core microservice for debugging purposes:

1. A system admin requests logs from core.

2. The admin receives the list of past events.

2.1.2   Supplementary Specification

Besides the functionalities specified in the previous section, the following non-functional requirements were pinned down.

Performance   Each instance of the data lake core should be able to respond up to 1000 simultaneous requests, which is approximated from the number of USTH researchers and students, every second.

Supportability   The data lake core and its dependencies must be able to run on common operating system (OS), including, but not limited to, GNU/Linux, Windows and macOS. While the microservice is likely to be deployed on GNU/Linux, it is uncertain if it will be the future maintainers' OS of choice for development. Furthermore, languages used for implementation and depended systems should be either familiar or easy to learn.

Licensing   The resulting software must be released under a copyleft license, in order to persist digital freedom in scientific research and education.

2.2 Design

2.2.1 Architecture

With consultancy from Dr Trần Giang Sơn, and collaboration with other interns[*], it was decided that in the data lake, the core will be communicating with the query engine and public API, which in turn take requests from other data lake services, e.g. crawler and dashboard, and external clients and transform them to comply with the core API. This is illustrated in figure 2.2.

On the other side, the core service encapsulates distributed file system (DFS) and database management system (DBMS) and provide a consistent interface for those storages. Therefore, the data lake core must also include clients to talk to these outside systems. For the content extraction use case, we added an *extractor* component reading from the file system (FS) and DB. The result will be either sent directly through the core API or cached in the DB.

2.2.2 Technology Choices

In a perfect world, choices of technology would be made following all other design decisions. However, existing technologies all have limitations (or at least trade-offs) that we need to be aware of to best decide on low-level details.

Programming Languages

Distributed File System

Database Management System

Logging

2.2.3 Interface

2.2.4 Database Schema

[*]Lê Như Chu Hiệp, Nguyễn Phương Thảo, Nguyễn An Thiết, Trần Minh Hiếu and Nguyễn Quốc Thông

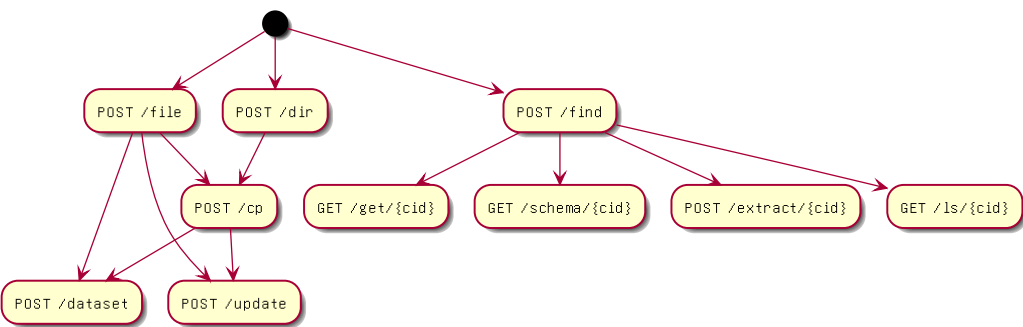Figure 2.2: Data lake overall architecture with focus on core's components

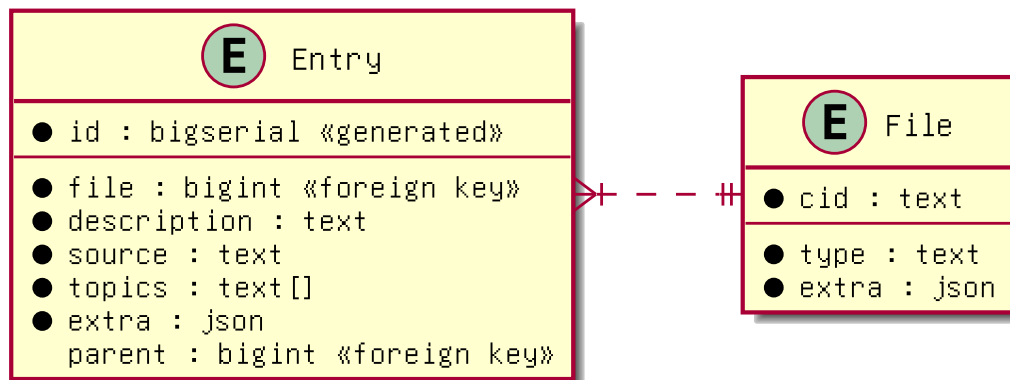Figure 2.3: Core HTTP API endpoints in a common order of access

Figure 2.4: Database schema

2.2.5 Query Abstract Syntax Tree

2.3 Implementation

2.3.1 Error Handling

2.3.2 I/O Handling

2.3.3 Query Transformations

2.3.4 Concurrency

2.4 Quality Assurance

3

Results and Discussion

3.1   Results

3.2   Discussion

4

Conclusion

A
Acknowledgment

This work would not have been possible without open access, Sci-Hub and Library Genesis. In addition, Wikipedia was a great index and summary for books and articles of certain topics.

Moreover, I would like to thank my colleagues and supervisors for their feedbacks, collaborations and encouragements. My family and friends (both online and in-real-life) were also incredibly supportive.

$\mathcal{B}$ Bibliography

[1] *Mission and Vision.* University of Science and Technology of Hanoi. `https://usth.edu.vn/en/abouts/Mission-et-Vision.html`.

[2] Huang Fang. "Managing Data Lakes in Big Data Era: What's a data lake and why has it became popular in data management ecosystem". *2015 IEEE International Conference on Cyber Technology in Automation, Control, and Intelligent Systems*, pp. 820–824. IEEE, 2015. doi:10.1109/CYBER.2015.7288049

[3] Chris Richardson. "1.4.1 Scale cube and microservices". *Microservice Patterns.* Manning Publications, 2018. ISBN 9781617294549.

[4] Brian Göetz, Tim Peierls, Joshua Bloch, Joseph Bowbeer and David Holmes. "3.4. Immutability". *Java Concurrency in Practice.* Addison Wesley Professional, 2006. ISBN 9780321349606.

[5] Alan Mark Davis. "Operational Prototyping: A new Development Approach". *IEEE Software*, vol. 9, no. 5, pp. 70–78. IEEE, 1992. doi:10.1109/52.156899.

C

Terms and Acronyms

Glossary

content a file or a directory. 2–5, 13

predicate an expression evaluating to either true or false. 3–5

Acronyms

API application programming interface. 1, 2, 6

CID content identifier. 2, 4, 5

DB database. 4, 6

DBMS database management system. 6

DFS distributed file system. 6

FS file system. 6

GNU GNU's not Unix. 5, 13

ID identifier. 2, 4

OS operating system. 5

USTH University of Science and Technology of Hanoi. 1, 2, 5, 12